

Find 10 different model providers and their models released in last 12 months

Model Provider	Model	Release Date	Cut-off Date	Cost of Tokens	Context Window	Type	When to Use
Open AI	GPT-4o	May-24	Apr-24	~50% cheaper than GPT-4	128K	Frontier multimodal	Real-time audio, vision, text tasks
Open AI	GPT-4o mini	Jul-24	Jun-24	97% cheaper than GPT-4	128K	Efficient	Everyday chat, lightweight apps
Anthropic	Claude 3 Opus	Mar-24	Feb-24	\$15 in / \$75 out per 1M tokens	200K	Frontier reasoning	Complex analysis, enterprise apps
Anthropic	Claude 3 Haiku	Mar-24	Feb-24	\$1 in / \$5 out per 1M tokens	200K	Fast/cheap	Quick responses, customer support
Google	Gemini 1.5 Pro	Feb-24	Jan-24	Proprietary pricing	1M	Frontier multimodal	Long-context reasoning, multimodal tasks
Google	Gemma 2 (9B/27B)	Jun-24	May-24	Free (open weights)	8K	Open-source	Research, fine-tuning
Meta	Llama 3.1 (8B–405B)	Jul-24	Jun-24	Free (open weights)	128K	Open-source	Tool use, large-context reasoning
Mistral	Mistral Large 2	Jul-24	Jun-24	Proprietary pricing	128K	Frontier	Competitive with GPT-4, enterprise
Microsoft	Phi-3 (mini/medium)	Apr-24	Mar-24	Free (open weights)	4K–128K	Small efficient	Mobile/edge deployment
Alibaba	Qwen 2 (0.5B–72B)	Jun-24	May-24	Free (open weights)	128K	Multilingual	Coding, multilingual tasks
Databricks	DBRX	Mar-24	Feb-24	Free (open weights)	32K	Enterprise MoE	Data/enterprise workloads
Cohere	Command R+	Apr-24	Mar-24	Free (open weights)	128K	Retrieval-augmented	Knowledge-heavy apps
NVIDIA	Nemotron 3 Ultra	Jun-25	May-25	\$0.60 in / \$2.60 out per 1M tokens	262K	Reasoning	High-performance enterprise AI